

# AI-Powered Task Scheduling Agent - AI Agents & Operations Documentation (Project Review Edition)

Project name: AI-Powered Task Scheduling Agent

System version: v2.1 (README "Latest Features"), API title indicates v2.0 (Week 4)

Backend: FastAPI + MongoDB + Firebase Auth + Socket.IO

AI runtime: Local Ollama models configured in settings (ollama\_model, ollama\_chat\_model) via backend/app/config.py

---

## 1) Executive Summary (Business + Technical)

This system embeds multiple task-focused AI agents into student/teacher workflows to reduce manual effort and increase planning quality. The AI agents are not separate microservices; they are internal modules invoked from FastAPI routers and services, primarily through a shared Ollama integration layer (backend/app/services/ollama\_service.py).

Business outcomes (as documented):

- Grading: ~70% faster
- Class monitoring: ~83% faster
- Task creation: ~90% faster

Technical approach:

- Most agents use prompted LLM inference (Ollama) with JSON-constrained outputs where structured data is required.
- Several agents include algorithmic fallbacks when AI fails (schedule generator fallbacks, JSON parsing fallbacks).

---

## 2) System AI Architecture (How Agents Fit Together)

### 2.1 High-level component flow

```
flowchart LR
    UI[React Frontend] --> API[FastAPI API]
    API --> DB[(MongoDB)]
    API --> WS[Socket.IO WebSocket]
    API --> AI[Ollama Service Layer]
    AI --> OLLAMA[(Local Ollama Runtime)]
    API --> GC[Google Calendar API]
    API --> OCR[Document Processor / OCR]

    subgraph Agents[AI Agents (Internal Modules)]
        A1[Task Analysis Agent]
        A2[Study Planner Agent]
        A3[Chat Assistant Agent]
        A4[Document Analysis Agent]
        A5[Resource Summarizer Agent]
        A6[Flashcard Agent]
        A7[Stress Recommendations Agent]
        A8[Extension Review Agent]
        A9[Grading Explanation Agent]
    end
```

```
API --> Agents  
Agents --> AI  
OCR --> Agents  
GC --> API  
WS --> UI
```

## 2.2 Shared AI integration layer

The backend centralizes model calls in `backend/app/services/ollama_service.py`:

- `generate_ai_response(prompt, json_mode=False)` uses `ollama.generate` for single-turn generation.
- `generate_json_response(prompt)` uses Ollama's JSON mode (`format='json'`) for better structured output.
- `generate_chat_response(user_message, system_prompt, chat_history, document_content)` uses `ollama.chat` for multi-turn, context-rich interactions.

---

## 3) AI Agent Catalog (All Implemented Agents)

### Conventions used in this section

- Agent Name + Version: The repository does not define per-agent semantic versions. Versions below reflect module maturity and the project "Week X" milestone notes. Where explicit versioning is unavailable, it is noted as "internal module (no explicit version tag)".
- I/O Interfaces: Listed as API endpoints and/or Python functions actually called.
- Dependencies: Python packages + internal modules + DB collections and integration APIs.

---

### 3.1 Task Complexity & Subtask Generation Agent (Week 1 core)

Official agent name: Task Complexity & Subtask Agent

Version: internal module (no explicit tag); used by Tasks API

#### Primary function & responsibilities

- Estimate complexity (1-10), estimated hours, priority, and suggested deadline offset
- Generate 3-5 actionable subtasks to accelerate task breakdown

#### Technical specifications

- Architecture: Prompted LLM ? JSON parsing with fallback defaults
- Algorithms: LLM-based estimation (no deterministic scoring)
- Structured output enforcement: "Respond ONLY in JSON" prompt + best-effort JSON extraction

#### Input/Output interfaces

- Python functions:
  - `backend/app/services/ai_task_service.py::analyze_task_complexity(title, description)`
  - `backend/app/services/ai_task_service.py::generate_subtasks(title, description)`
- API integration:
  - POST /api/tasks/ calls both functions (`backend/app/routers/tasks.py`)

#### Dependencies and integration points

- AI: `backend/app/services/ollama_service.py::generate_ai_response`
- MongoDB: `tasks_collection` (task insertion)

- WebSocket: broadcasts task updates via backend/app/websocket/broadcaster.py
- Optional Calendar: background sync via backend/app/services/google\_calendar\_service.py if enabled

Concrete example

- User action: Student creates a task via POST /api/tasks/
  - Agent input: title + description
  - Expected AI output shape:
- ```
```json
{"complexity": 6, "hours": 5, "deadline_days": 7, "priority": "high"}
```

```

- System result: Task stored with complexity\_score, estimated\_hours, ai\_suggested\_deadline, and AI-generated subtasks.

---

### 3.2 Smart Study Planner Agent (Week 3)

Official agent name: Smart Study Planner Agent

Version: internal module (Week 3)

Primary function & responsibilities

- Generate a daily (or weekly) study schedule with:
- Deadline urgency ordering
- Complexity/stress balancing
- Break insertion (short/long)
- Session typing (pomodoro深深呼吸/short\_burst)

Technical specifications

- Hybrid architecture:
  - 1) Deterministic scoring + slot calculation
  - 2) LLM attempts a schedule first
  - 3) Deterministic greedy fallback if LLM fails
- Algorithms:
  - Weighted scoring model: backend/app/services/ai\_scheduling\_service.py::score\_tasks
  - LLM schedule proposal (JSON): backend/app/services/ai\_scheduling\_service.py::ai\_generate\_schedule
  - Greedy fallback: backend/app/services/ai\_scheduling\_service.py::generate\_fallback\_schedule
  - Orchestrator: backend/app/services/ai\_scheduling\_service.py::generate\_study\_schedule

Input/Output interfaces

- API endpoints:
  - POST /api/study-planner/generate (daily): backend/app/routers/study\_planner.py
  - POST /api/study-planner/generate-week (weekly): backend/app/routers/study\_planner.py
- Inputs:
  - target\_date, regenerate
  - user preferences (study hours, session length, breaks, complexity pattern)
  - latest stress score (from stress logs)
- Output:
  - JSON schedule doc with study\_blocks, break\_blocks, total\_study\_hours, ai\_reasoning

Dependencies and integration points

- MongoDB:
- reads tasks\_collection (active tasks)
- reads stress\_logs\_collection (latest stress)
- writes study\_plans\_collection
- AI: backend/app/services/ollama\_service.py::generate\_ai\_response

- Preferences: user\_preferences\_collection

Concrete example

- User action: Student clicks "Generate Today's Plan"
- System call: POST /api/study-planner/generate
- Typical output shape:

```
```json
{
  "study_blocks": [{"task_title": "Math Assignment", "start_time": "10:00", "end_time": "10:25", "duration_minutes": 25, "session_type": "pomodoro"}],
  "break_blocks": [{"start_time": "10:25", "end_time": "10:30", "duration_minutes": 5, "break_type": "short"}],
  "total_study_hours": 3.5,
  "ai_reasoning": "High urgency tasks first; shorter sessions due to stress."
}
```

---

### 3.3 Context-Aware AI Chat Assistant Agent (Week 4+ Enhanced)

Official agent name: Context-Aware AI Assistant (Chat)

Version: internal module; "basic" and "enhanced" endpoints implemented

Primary function & responsibilities

- Provide conversational assistance about:
- tasks, deadlines, priorities
- study plans and wellbeing signals (enhanced)
- slash-command execution (enhanced)
- document Q&A (enhanced)

Technical specifications

- Basic mode: prompt string composition + ollama.generate (single-turn)
- Enhanced mode: system prompt + scoped context + recent chat history + optional document content ? ollama.chat
- Context aggregation: multi-collection context builder: backend/app/services/user\_context\_service.py

Input/Output interfaces

- API endpoints:
- POST /api/chat/ai (basic): backend/app/routers/chat.py
- POST /api/chat/ai/enhanced (enhanced, file + scope + commands): backend/app/routers/chat.py
- POST /api/chat/ai/command (execute slash command): backend/app/routers/chat.py
- GET /api/chat/ai/context (context preview): backend/app/routers/chat.py

Dependencies and integration points

- AI:
  - basic: backend/app/services/ollama\_service.py::generate\_ai\_response
  - enhanced: backend/app/services/ollama\_service.py::generate\_chat\_response
- Context:
  - backend/app/services/user\_context\_service.py::get\_full\_user\_context
  - backend/app/services/user\_context\_service.py::format\_context\_for\_ai
- Slash commands:
  - backend/app/services/command\_parser.py
- Storage:
  - AI conversations stored in chat\_history\_collection

## Concrete examples

### 1) Task question (enhanced)

- Input: "What are my overdue tasks?"

- Behavior: the system injects overdue/pending task context into the system prompt and answers with task names and deadlines.

### 2) Command execution

- Input: /task list

- Behavior: command is executed server-side; result is stored as an assistant message and returned.

---

## **3.4 Document Analysis Agent (Upload ? Extract ? Answer)**

Official agent name: Document Analysis Assistant

Version: internal module

### Primary function & responsibilities

- Extract text from user-uploaded documents (PDF/DOCX/text/code/images via OCR)

- Provide AI analysis or Q&A over extracted content

### Technical specifications

- Extraction pipeline: backend/app/services/document\_processor.py

- PDF: pypdf

- DOCX: python-docx

- Images: Pillow + pytesseract OCR

- Analysis:

- backend/app/services/ollama\_service.py::analyze\_document\_with\_ai

- Enhanced chat integrates extraction directly in backend/app/routers/chat.py

### Input/Output interfaces

- Input:

- file upload to POST /api/chat/ai/enhanced (multipart form file)

- Output:

- AI response plus document\_metadata (type, filename, character count)

### Dependencies and integration points

- Settings:

- ai\_max\_document\_size, ai\_max\_context\_length from backend/app/config.py

- OCR:

- Windows Tesseract path detection in backend/app/services/document\_processor.py

### Concrete example

- User uploads a PDF lecture note and asks: "Summarize key points and give practice questions."

- System extracts PDF text, truncates if needed, injects into AI chat context, returns a structured explanation.

---

## **3.5 Resource Summarization & Auto-Tagging Agent (Resource Library)**

Official agent name: Resource Summarizer Agent

Version: internal module (Week 1)

### Primary function & responsibilities

- Generate concise summaries of notes/documents and extract key points (used as suggested tags/key

concepts).

#### Technical specifications

- LLM prompt with bounded preview (content[:3000])
- Best-effort JSON extraction + fallback summary

#### Input/Output interfaces

- Function: backend/app/routers/resources.py::generate\_resource\_summary(content, title)
- Usage:
- file uploads and note creation in backend/app/routers/resources.py

#### Dependencies and integration points

- AI: backend/app/services/ollama\_service.py::generate\_ai\_response
- MongoDB: resources\_collection

#### Concrete example

- Teacher uploads "Syllabus.pdf" ? system stores file + ai\_summary + ai\_key\_points ? UI shows summary and suggested tags.

---

### **3.6 Flashcard Generation Agent (Resources + Chat Command)**

Official agent name: Flashcard Generator Agent

Version: internal module (Week 1)

#### Primary function & responsibilities

- Generate 8-12 study flashcards from a note/document summary/content.

#### Technical specifications

- Uses JSON-constrained generation mode:  
backend/app/services/ollama\_service.py::generate\_json\_response
- Validates card schema and filters invalid entries.

#### Input/Output interfaces

- Endpoint: POST /api/resources/{resource\_id}/flashcards (backend/app/routers/resources.py)
- Core generator: backend/app/routers/resources.py::generate\_flashcards\_ai
- Chat command: /flashcard generate <topic> (implemented in backend/app/services/command\_parser.py)

#### Dependencies and integration points

- AI: JSON output mode for higher success rate
- MongoDB: resources\_collection
- Extraction fallback:
- flashcard endpoint tries to read existing PDFs/text files if stored content is missing

#### Concrete example

- Student clicks "Generate Flashcards" on a note; system generates and persists flashcards and renders them in the flashcard viewer.

---

### **3.7 Stress Recommendations Agent (Wellbeing)**

Official agent name: Stress Recommendation Agent

Version: internal module (Week 1)

#### Primary function & responsibilities

- Provide actionable recommendations to reduce stress based on workload and deadlines.

#### Technical specifications

- Deterministic stress scoring (0-10) based on:
  - upcoming deadline pressure
  - average complexity
  - deadline overlap
  - historical pattern from recent logs
- LLM recommendations: attempts JSON array parsing, otherwise uses fallback recommendations

#### Input/Output interfaces

- Endpoint: GET /api/stress/current (backend/app/routers/stress.py)
- Persists results to stress\_logs collection.

#### Dependencies and integration points

- MongoDB:
- reads tasks\_collection
- reads/writes stress\_logs
- AI: backend/app/services/ollama\_service.py::generate\_ai\_response

#### Concrete example

- Student has multiple near-term deadlines ? objective stress increases ? agent recommends prioritization and timeboxing actions.

---

### 3.8 Extension Request Review Agent (Teacher workflow support)

Official agent name: Extension Review Agent

Version: internal module

#### Primary function & responsibilities

- Analyze a student's deadline extension request and recommend approve/deny/conditional with confidence and reasoning.

#### Technical specifications

- LLM prompt forces JSON response schema; parses JSON with fallback
- Confidence clamped to [0.0, 1.0]

#### Input/Output interfaces

- Function: backend/app/services/ai\_extension\_service.py::analyze\_extension\_request(...)
- Endpoint: POST /api/extensions/ (backend/app/routers/extensions.py) stores ai\_recommendation, ai\_reasoning, ai\_confidence\_score.

#### Dependencies and integration points

- MongoDB: extension\_requests\_collection, tasks\_collection, notifications\_collection
- WebSocket: notification broadcast to teacher via backend/app/websocket/broadcaster.py
- AI: backend/app/services/ollama\_service.py::generate\_ai\_response

#### Concrete example

- Student requests +7 days due to overlapping deadlines ? agent suggests "conditional" with a suggested deadline ? teacher approves/denies via review endpoint.

---

### **3.9 AI Grading Explanation Agent (Implemented, limited integration)**

Official agent name: AI Grading Assistant (Explanation)

Version: internal module (Week 2 feature)

Primary function & responsibilities

- Turn performance metrics into a teacher-friendly, constructive grading explanation:
- reasoning
- strengths, weaknesses
- improvement suggestions
- encouragement

Technical specifications

- LLM prompt asks for JSON response; parses JSON; fallback response if invalid
- Includes deterministic grade calculation helper: calculate\_grade\_from\_performance

Input/Output interfaces

- Functions:

- backend/app/services/ai\_grading\_service.py::generate\_grading\_explanation(...)
- backend/app/services/ai\_grading\_service.py::calculate\_grade\_from\_performance(...)
- Note: current teacher review endpoints in backend/app/routers/grading.py focus on task review and feedback submission; direct invocation of ai\_grading\_service is not shown in the current router.

Dependencies and integration points

- AI: backend/app/services/ollama\_service.py::generate\_ai\_response

Concrete example

- Intended usage: compile submission metrics + call agent ? return structured feedback JSON suitable for a grading dashboard UI.

---

## **4) Implementation Details**

### **4.1 Development timeline and version history**

- Week 1: Student AI features (stress meter, focus mode, resource AI summary + flashcards)
- Week 2: Teacher efficiency tools (grading assistant module, class analytics, bulk tasks)
- Week 3: Smart study planner
- Week 4: Google Calendar integration
- v2.1 (Jan 2026): UI redesign + WebSocket features + USN support + resource library enhancements

### **4.2 Performance metrics and benchmarks**

Business-level metrics available in-repo

- Teacher time savings and efficiency improvements are documented in README.md.

System-level benchmarks (current state)

- The backend does not include explicit latency/throughput benchmarking instrumentation.
- Performance drivers:
  - LLM inference time (model size, hardware, prompt length)
  - MongoDB query performance (indexes configured in backend/app/db\_config.py)
  - WebSocket fanout volume for broadcasts

Recommended benchmark plan for review

- p50/p95 latency per AI endpoint (/chat/ai/enhanced, /study-planner/generate, /resources/{id}/flashcards)
- Prompt size vs latency curve (document uploads + large contexts)
- Fallback rate (% JSON parsing failures leading to fallback behavior)
- Error rate categories (Ollama unreachable, OCR missing, malformed docs)

### 4.3 Error handling and recovery mechanisms

- Central AI error pattern: generate\_ai\_response returns strings beginning with "AI Error:" on failure.
- Structured output fallback: agents attempt to parse JSON by slicing {...} or [...]; on failure they return safe defaults.
- Task analysis defaults in backend/app/services/ai\_task\_service.py
- Scheduling fallback algorithm in backend/app/services/ai\_scheduling\_service.py
- Flashcards return structured errors and validate schema in backend/app/routers/resources.py
- Document extraction resilience: if optional libraries are missing, the extraction returns explicit error text; unsupported types are rejected by allow-list.

### 4.4 Security considerations

Authentication and authorization

- REST: Firebase token verification in routers (example: backend/app/routers/tasks.py)
- WebSocket: token required on connect and validated in backend/app/websocket/events.py

Secrets management

- Required env vars: MONGODB\_URI, FIREBASE\_CREDENTIALS\_PATH, SECRET\_KEY in backend/app/config.py
- Calendar token encryption key: optional but strongly recommended (CALENDAR\_ENCRYPTION\_KEY)

Sensitive token storage

- Google tokens are encrypted with Fernet before storage in MongoDB (backend/app/services/google\_calendar\_service.py).
- Operational caveat: if CALENDAR\_ENCRYPTION\_KEY is not set, an ephemeral runtime key is generated, which can make stored tokens undecryptable after restart. Workaround: set a stable CALENDAR\_ENCRYPTION\_KEY.

File upload safety

- AI chat document upload:
- extension allow-list via is\_supported\_file
- size limit via ai\_max\_document\_size
- filename sanitization and safe storage under uploads/ai\_docs/...
- Task attachment upload:
- allow-listed extensions + max size in backend/app/routers/tasks.py

---

## 5) Operational Documentation

### 5.1 Deployment configuration

Local/dev startup

- Windows starter script: start.bat
- Backend: uvicorn app.main:app --reload --host 0.0.0.0 --port 8000

- Frontend: npm run dev
- Combined REST + Socket.IO ASGI app is wired in backend/app/main.py.

Key runtime dependencies

- Ollama local runtime
- MongoDB instance
- Firebase service account JSON
- Optional: Google OAuth credentials for calendar sync

## 5.2 Monitoring and maintenance procedures

- Logging writes:
  - logs/app.log (DEBUG+)
  - logs/error.log (ERROR+)
  - console INFO logs
  - configured in backend/app/utils/logger.py
- Operational checks:
  - GET /health endpoint in backend/app/main.py
  - AI connectivity sanity check: backend/app/services/ollama\_service.py::test\_ollama\_connection

## 5.3 Scaling capabilities

- API scaling: FastAPI can scale horizontally behind a reverse proxy (stateless requests).
- WebSocket scaling: Socket.IO generally requires sticky sessions or a shared broker to scale horizontally. Current configuration appears single-process (no broker configured).
- AI scaling: local Ollama runtime is typically the bottleneck; scale vertically (GPU/CPU/RAM) or isolate Ollama onto a dedicated host.

## 5.4 Known limitations and workarounds

- Calendar encryption key stability: set CALENDAR\_ENCRYPTION\_KEY to persist decryptability across restarts.
- Per-agent versioning: modules do not have explicit semantic versions; add explicit version tags if formal auditability is required.
- AI grading integration gap: grading explanation module exists but is not visibly invoked by grading endpoints; integrate into teacher review flow if required.
- Benchmark instrumentation absent: business metrics exist, but p95 latency/fallback rates/error budgets are not measured in code; add telemetry if production review requires SLAs.

---

## Appendix A) Quick Reference: AI-Related Endpoints (Backend)

- Tasks (AI analysis + subtasks): POST /api/tasks/ (backend/app/routers/tasks.py)
- Study planner: POST /api/study-planner/generate (backend/app/routers/study\_planner.py)
- Stress recommendations: GET /api/stress/current (backend/app/routers/stress.py)
- Resource summarization: note/file endpoints in backend/app/routers/resources.py
- Flashcards: POST /api/resources/{resource\_id}/flashcards (backend/app/routers/resources.py)
- Extensions (AI recommendation): POST /api/extensions/ (backend/app/routers/extensions.py)
- AI chat:
- Basic: POST /api/chat/ai
- Enhanced: POST /api/chat/ai/enhanced (backend/app/routers/chat.py)